

Node.js

Mendel Rosenblum

Threads

versus

Events

```
request = readRequest(socket);  
reply = processRequest(request);  
sendReply(socket, reply);
```

Implementation:

Thread switching (i.e. blocking) and a scheduler

```
readRequest(socket, function(request) {  
    processRequest(request,  
        function (reply) {  
            sendReply(socket, reply);  
        });  
});
```

Implementation:

Event queue processing

Event queue

- Inner loop

```
while (true) {  
    if (!eventQueue.isEmpty()) {  
        eventQueue.pop().call();  
    }  
}
```

- Never wait/block in event handler . Example `readRequest(socket);`

1. `launchReadRequest(socket);` // Returns immediately
2. When read finishes: `eventQueue.push(readDoneEventHandler);`

Node.js

- Take a JavaScript engine from a browser (Chrome's V8 JavaScript Engine)
 - Get same JavaScript on both browser and server
 - Don't need the DOM on the server
- Add events and an event queue
 - Everything runs as a call from the event loop (already had one for browser events)
- Make event interface to all OS operations
 - Wrap all the OS blocking calls (file and socket/network io)
 - Add some data handle support
- Add a proper module system
 - Each module gets its own scope (not everything in window)

Example: Node.js reading a file

```
var fs = require("fs"); // require is a Node module call
                        // fs object wraps OS sync file system calls

// OS read() is synchronous but Node's fs.readFile is asynchronous
fs.readFile("smallFile", readDoneCallback); // Start read

function readDoneCallback(error, dataBuffer) {
    // Node callback convention: First argument is JavaScript Error object
    // dataBuffer is a special Node Buffer object
    if (!error) {
        console.log("smallFile contents", dataBuffer.toString());
    }
}
```

Node Modules

- Import using `require()`
 - System module: `require("fs");` // Looks in `node_module` directories
 - From a file: `require("./XXX.js");` // Reads specified file
 - From a directory: `require("./myModule");` // Reads `myModule/index.js`
- Module files have a private scope
 - Can declare variables that would be global in the browser
 - `Require` returns what is assigned to `module.exports`

```
var notGlobal;  
function func1() {}  
function func2() {}  
module.exports = {func1: func1, func2: func2};
```

Node modules

- Many standard Node modules
 - File system, process access, networking, timers, devices, crypto, etc.
- Huge library of modules (npm)
 - Do pretty much anything you want
- We use:
 - Express - Fast, unopinionated, minimalist web framework (speak HTTP)
 - Mongoose - Elegant mongodb object modeling (speak to the database)

Node Buffer class

- Manipulating lots of binary data wasn't a strength of the JavaScript engine
 - Unfortunately that is what web servers do: DBMS ⇔ Web Server ⇔ Browser
- Node add a Buffer class - Optimized for storing and operating on binary data
 - Interface looks an array of bytes (like the OS system calls use)
 - Memory is allocated outside of the V8 heap
- Used by the wrapped OS I/O calls (fs, net, ...)
- Optimized sharing with pointers rather than always copying
 - `buffer.copy()`
 - For example: `fs.readFile` to `socket.write`

Buffer operations

- Supports operations for picking values out or updating them
 - Can peek at some values and send the bulk of it on its way
- Can convert a buffer or parts of a buffer to a JavaScript string
 - `buf.toString("utf8");` // Convert to UTF8 - commonly used on the web
 - `buf.toString("hex");` // Convert to hex encoding (2 digits per byte)
 - `buf.toString("base64");` // Convert to base64 encoding

Example: Node.js reading a file (redux)

```
var fs = require("fs");  
// fs has 81 properties readFile, writeFile, most OS calls, etc.  
fs.readFile("smallFile", readDoneCallback); // Start read  
// Read has been launched - JavaScript execution continues  
// Node.js exits when no callbacks are outstanding  
  
function readDoneCallback(error, dataBuffer) {  
    // console.log(dataBuffer) prints <Buffer 66 73 20 3d 20 72 65 71 ...  
    if (!error) {  
        console.log("smallFile contents", dataBuffer.toString());  
    }  
}
```

Programming with Events/Callbacks

- Key difference
 - Threads: Blocking/waiting is transparent
 - Events: Blocking/waiting requires callback

- Mental model

If code doesn't block: Same as thread programming

If code does block (or needs to block): Need to setup callback

Often what was a return statement becomes a function call

Example: Three step process

Threads

```
r1 = step1();
console.log('step1 done', r1);
r2 = step2(r1);
console.log('step2 done', r2);
r3 = step3(r2);
console.log('step3 done', r3);
console.log('All Done!');
```

Works for **non-blocking**
calls in both styles

Callbacks

```
step1(function(r1) {
  console.log('step1 done', r1);
  step2(r1, function (r2) {
    console.log('step2 done', r2);
    step3(r2, function (r3) {
      console.log('step3 done',r3);
    });
  });
}):
console.log('All Done!'); // Wrong!
```

Example: Three step process

Threads

```
r1 = step1();  
console.log('step1 done', r1);  
r2 = step2(r1);  
console.log('step2 done', r2);  
r3 = step3(r2);  
console.log('step3 done', r3);  
console.log('All Done!');
```

Callbacks

```
step1(function(r1) {  
    console.log('step1 done', r1);  
    step2(r1, function (r2) {  
        console.log('step2 done', r2);  
        step3(r2, function (r3) {  
            console.log('step3 done', r3);  
            console.log('All Done!');  
        });  
    });  
}):
```

Listener/emitter pattern

- When programing with events (rather than threads) a listener/emitter pattern is used.
- Listener - Function to be called when the event is signaled
 - Should be familiar from DOM programming (`addEventListener`)
- Emitter - Signal that an event has occurred
 - Emit an event cause all the listener functions to be called

Node: EventEmitter = require('events');

- Listen with `on()` and signal with `emit()`

```
myEmitter.on('myEvent', function(param1, param2) {  
    console.log('myEvent occurred with ' + param1 + ' and ' + param2 + '!');  
});  
myEmitter.emit('myEvent', 'arg1', 'arg2');
```

- On emit call listeners are called synchronously and in the order the listeners were registered
- If no listener then `emit()` is a nop

Typical EventEmitter patterns

- Have multiple different events for different state or actions

```
myEmitter.on('conditionA', doConditionA);
```

```
myEmitter.on('conditionB', doConditionB);
```

```
myEmitter.on('conditionC', doConditionC);
```

```
myEmitter.on('error', handleErrorCondition);
```

- Handling 'error' is important - Node exits if not caught!

```
myEmitter.emit('error', new Error('Ouch!'));
```


Streams

- Build modules that produce and/or consume streams of data
- A popular way of structuring servers
 - Network interface \Leftrightarrow TCP/IP protocol processing \Leftrightarrow HTTP protocol processing \Leftrightarrow your code
- Can build connected streams dynamically
 - Add modules on stream: E.g. `stream.push(Encryption)`
Network interface \Leftrightarrow TCP/IP protocol processing \Leftrightarrow **Encryption** \Leftrightarrow HTTP processing
- Node's APIs heavily uses streams
 - Readable streams (e.g. `fs.createReadStream`)
 - Writable stream (e.g. `fs.createWriteStream`)
 - Duplex stream (e.g. `net.createConnection`)
 - Transform stream (e.g. `zlib`, `crypto`)

Readable streams - File reading using streams

```
var readableStreamEvent = fs.createReadStream("smallFile");

readableStreamEvent.on('data', function (chunkBuffer) { // Could be called multiple times
  console.log('got chunk of', chunkBuffer.length, 'bytes');
});

readableStreamEvent.on('end', function() {
  // Called after all chunks read
  console.log('got all the data');
});

readableStreamEvent.on('error', function (err) {
  console.error('got error', err);
});
```

Writable streams - File writing using streams

```
var writableStreamEvent = fs.createWriteStream('outputFile');
```

```
writableStreamEvent.on('finish', function () {  
    console.log('file has been written!');  
});
```

```
writableStreamEvent.write('Hello world!\n');
```

```
writableStreamEvent.end();
```

Digression: Socket setup for TCP connections

Client (Browser)

```
sock = socket(AF_INET, SOCK_STREAM, 0);  
  
connect(sock, &serverAddr,  
        sizeof(serverAddr));  
  
write(sock, "Hi!", 3);  
  
read(sock, buf, 3)
```

Server (Web Server)

```
lfd = socket(AF_INET, SOCK_STREAM, 0);  
  
bind(lfd, &serverSaddr, sizeof(serveraddr));  
  
listen(lfd, 5);  
  
sock = accept(lfd, &clientaddr, &clientlen);  
  
read(sock, buf, 3);  
  
write(sock, buf, 3)
```

- TCP/IP socket connection is a reliable, in-order byte stream
 - Note: reads can return data in different chunks that sent

TCP Networking on Node.js

- Node net module wraps OS's network routines
- Includes higher level functionality like:

```
var net = require('net');
```

```
net.createServer(processTCPconnection).listen(4000);
```

Creates a socket, binds port 4000, and listens for connections

Calls function `processTCPconnection` on each TCP connection

Example: A chat server

```
var clients = []; // List of connected clients
function processTCPconnection(socket) {
  clients.push(socket); // Add this client to our connected list
  socket.on('data', function (data) {
    broadcast( "> " + data, socket); // Send received data to all
  });
  socket.on('end', function () {
    clients.splice(clients.indexOf(socket), 1); // remove socket
  });
}
```

Chat Server: broadcast

```
// Send message to all clients
function broadcast(message, sender) {
  clients.forEach(function (client) {
    if (client === sender) return; // Don't send it to sender
    client.write(message);
  });
}
```

- Our chat server implementation is done!

Putting it together: A real simple file server

```
net.createServer(function (socket) {  
  socket.on('data', function (fileName) {  
    fs.readFile(fileName.toString(), function (error, fileData) {  
      if (!error) {  
        socket.write(fileData); // Writing a Buffer  
      } else {  
        socket.write(error.message); // Writing a String  
      }  
      socket.end();  
    });  
  });  
}).listen(4000);
```

- Think about concurrency going on here

Example: Read three files

- Linux

```
read(open("f1"), buf1, f1Size);  
read(open("f2"), buf2, f2Size);  
read(open("f3"), buf3, f3Size);
```

- Node.js

```
fs.readFile("f1", function (error, data1) {  
    fs.readFile("f2", function (error, data2) {  
        fs.readFile("f3", function (error, data3) {  
            // Call Pyramid of Doom or Callback Hell  
        });  
    });  
});
```

Example: Read N files

```
var fileContents = {};  
['f1','f2','f3'].forEach(function (fileName) {  
    fs.readFile(fileName, function (error, dataBuffer) {  
        assert(!error);  
        fileContents[fileName] = dataBuffer;  
    });  
});
```

If we want to use `fileContents` how do we know when all reads are finished?

Recall: Can't wait in NodeJS

Example: Read N files

```
var fileContents = {};  
['f1','f2','f3'].forEach(function (fileName) {  
    fs.readFile(fileName, function (error, dataBuffer) {  
        assert(!error);  
        fileContents[fileName] = dataBuffer;  
        if (gotLastCallback()) allDoneCallback(fileContents);  
    });  
});
```

- Yuck!

Async module: `var async = require('async');`

- Solution: Write a function that turns waiting into a callback

```
var fileContents = {};  
  
async.each(['f1','f2','f3'], readIt, function (err) {  
    if (!err) console.log('Done'); // fileContents filled in  
    if (err) console.error('Got an error:', err.message);  
});  
  
function readIt(fileName, callback) {  
    fs.readFile(fileName, function (error, dataBuffer) {  
        fileContents[fileName] = dataBuffer;  
        callback(error);  
    });  
}
```

Node.JS - many useful built-in modules

- Buffer
- C/C++ Addons
- Child Processes
- Cluster
- Console
- Crypto
- Debugger
- DNS
- Errors
- Events
- File System
- Globals
- HTTP
- HTTPS
- Modules
- Net
- OS
- Path
- Process
- Punycode
- Query Strings
- Readline
- REPL
- Stream
- String Decoder
- Timers
- TLS/SSL
- TTY
- UDP/Datagram
- URL
- Utilities
- V8
- VM
- ZLIB